

Design Patterns for Semaphores

Kenneth A. Reek, Professor
Department of Computer Science
102 Lomb Memorial Drive
Rochester Institute of Technology
Rochester, NY 14623

kar@cs.rit.edu

Introduction

Synchronization is difficult!

- Synchronization using low-level techniques is more difficult
- But high-level techniques may not always be available
- Teaching high-level techniques does not convey the underlying concepts as well

Most textbooks describe semaphores

- Illustrate simple mutual exclusion
- Show solutions for a couple of classic problems, possibly including some “tricks”
- Do not describe general approaches

Readers/Writers

Rules:

- Any number of readers can access the file simultaneously if there are no writers
- Each writer must have exclusive access to the file

Rules to avoid indefinite postponement:

- When one or more writers is waiting to access the file, newly arriving readers are blocked
- When an active writer departs, waiting readers are unblocked in preference to unblocking another writer

System State

The state of this system includes four counts:

- Number of readers active and blocked
- Number of writers active and blocked

Might be many processes testing and changing state

- The state itself is a critical resource
- Accesses to it must be protected with mutual exclusion
- But you must not block yourself inside of the mutual exclusion!

Incorrect Solution

```
// READERS
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P();
    mutex.P();
    r_waiting -= 1;
}
r_active += 1;
mutex.V();
// ... read the file ...
mutex.P();
r_active -= 1;
if( r_active == 0 && w_waiting > 0 )
    writer.V();
mutex.V();
```

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P();
    mutex.P();
    w_waiting -= 1;
}
w_active += 1;
mutex.V();
// ... write the file ...
mutex.P();
w_active -= 1;
if( r_waiting > 0 )
    for( int i = 0; i < r_waiting; i += 1 )
        reader.V();
else if( w_waiting > 0 )
    writer.V();
mutex.V();
```

Incorrect Solution: What's Wrong

The problem is that unblocked processes must *reenter* the mutual exclusion

- Newly arriving processes might beat them to it
- When unblocked processes finally run, the system state may have changed—it may no longer be legal for them to run

Incorrect Solution: Example

A writer is accessing the file and several readers are blocked

r_active: 0 r_waiting: 3 w_active: 1 w_waiting: 0

```
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P(); ← 3 blocked
    mutex.P();
    r_waiting -= 1;
}
r_active += 1;
mutex.V();
// ... read the file ...
```

```
// ... write the file ...
mutex.P();
w_active -= 1;
if( r_waiting > 0 )
    for( int i = 0; i < r_waiting; i += 1 )
        reader.V();
else if( w_waiting > 0 )
    writer.V();
mutex.V();
```

1) The writer leaves, unblocking all the blocked readers

r_active: 0 r_waiting: 3 w_active: 0 w_waiting: 0

Incorrect Solution: Example (continued)

- 2) Before the readers get a quantum, a new writer arrives and “cuts in line”

r_active: 0 r_waiting: 3 w_active: 0 w_waiting: 0

```
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P(); ← 3 unblocked
    mutex.P();
    r_waiting -= 1;
}
r_active += 1;
mutex.V();
// ... read the file ...
```

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P();
    mutex.P();
    w_waiting -= 1;
}
w_active += 1;
mutex.V();
// ... write the file ...
```

r_active: 0 r_waiting: 3 w_active: 1 w_waiting: 0

Incorrect Solution: Example (continued)

- 3) The unblocked readers, one by one, resume and begin reading

r_active: 0 r_waiting: 3 w_active: 1 w_waiting: 0

```
// READERS
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P(); ← 3 unblocked processes resume
    mutex.P();
    r_waiting -= 1;
}
r_active += 1;
mutex.V();
// ... read the file ...
```

r_active: 3 r_waiting: 0 w_active: 1 w_waiting: 0

Incorrect Solution: Explanation

The problem is the delay between *unblocking* the processes and *updating the system state* to reflect the unblocking

- The unblocked processes might resume at any time
- You must consider them “running” *as soon as they are unblocked*
- Cannot depend on the unblocked processes themselves to update the system state

A Possible Fix

Change the “if” to a “while”

```
// READERS
mutex.P();
while( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P();
    mutex.P();
    r_waiting -= 1;
}
r_active += 1;
mutex.V();
// ... read the file ...
```

- Unfair to waiting processes
- Exacerbates indefinite postponement

A Better Fix

Update the system state when the unblocking is done

- The unblocking process does the updating on behalf of the unblocked process
- This is the “I’ll Do It For You” pattern

Implementing “I’ll Do It For You”

- 1) Identify the statements executed by unblocked processes
- 2) Change the code so that unblocked processes do not execute these statements

```
// READERS
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P();
    mutex.P();
    r_waiting -= 1;
}
r_active += 1;
mutex.V();
// ... read the file ...
```

```
// READERS
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P();
} else {
    r_active += 1;
    mutex.V();
}
```

Implementing “I’ll Do It For You” (continued)

- 3) Put the same functionality into the unblocking code

```
// ... write the file ...
mutex.P();
w_active -= 1;
if( r_waiting > 0 )
    for( int i = 0; i < r_waiting; i += 1 )
        reader.V();
else if( w_waiting > 0 )
    ...
```

```
// ... write the file ...
mutex.P();
w_active -= 1;
if( r_waiting > 0 ){
    r_active = r_waiting;
    while( r_waiting > 0 ){
        r_waiting -= 1;
        reader.V();
    }
} else if( w_waiting > 0 )
    ...
```

The unblocked processes do not reenter the mutual exclusion at all

The “I’ll Do It For You” Solution

```
// READERS
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P();
} else {
    r_active += 1;
    mutex.V();
}
// ... read the file ...
mutex.P();
r_active -= 1;
if( r_active == 0 && w_waiting > 0 ){
    w_waiting -= 1;
    w_active += 1;
    writer.V();
}
mutex.V();
```

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P();
} else {
    w_active += 1;
    mutex.V();
}
// ... write the file ...
mutex.P();
w_active -= 1;
if( r_waiting > 0 ){
    r_active = r_waiting;
    while( r_waiting > 0 ){
        r_waiting -= 1;
        reader.V();
    }
} else if( w_waiting > 0 ){
    w_waiting -= 1;
    w_active += 1;
    writer.V();
}
mutex.V();
```

*Rochester Institute of Technology
Department of Computer Science*

I'll Do It For You: Characteristics

The unblocked processes do not reenter the mutual exclusion

- Unblocked process cannot access any state variables
- You can unblock any number of processes
- Cohesion suffers

I'll Do It For You: Characteristics (continued)

- Still vulnerable to “cutting in line,” but much less likely
- 1) One process is reading the file, and a writer arrives

r_active: 1 r_waiting: 0 w_active: 0 w_waiting: 0

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P(); ← process loses its quantum here
                (very unlikely, but possible)
} else {
    w_active += 1;
    mutex.V();
}
// ... write the file ...
```

r_active: 1 r_waiting: 0 w_active: 0 w_waiting: 1

I'll Do It For You: Characteristics (continued)

2) The reader leaves and tries to unblock the writer

r_active: 1 r_waiting: 0 w_active: 0 w_waiting: 1

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P(); ← process paused here
} else {
    w_active += 1;
    mutex.V();
}
// ... write the file ...
```

```
// ... read the file ...
mutex.P();
r_active -= 1;
if( r_active == 0 && w_waiting > 0 ){
    w_waiting -= 1;
    w_active += 1;
    writer.V();
}
mutex.V();
```

r_active: 0 r_waiting: 0 w_active: 1 w_waiting: 0

I'll Do It For You: Characteristics (continued)

- 3) A new writer arrives before the original one resumes

r_active: 0 r_waiting: 0 w_active: 1 w_waiting: 0

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P(); ← process paused here
} else {
    w_active += 1;
    mutex.V();
}
// ... write the file ...
```

r_active: 0 r_waiting: 0 w_active: 1 w_waiting: 1

Only way to guarantee ordering is to block each process on a separate semaphore

Alternate Approach

The original problem is caused by having to *reenter* the mutual exclusion

- The unblocking process leaves without opening the mutual exclusion
- The unblocked process takes over the mutual exclusion
- This is the **Pass the Baton** pattern

Like runners in a relay race passing the baton to one another

- Occupancy of the mutual exclusion is our “baton”

Implementing “Pass the Baton”

- 1) Find where unblocked processes reenter the mutual exclusion
- 2) Eliminate these P operations

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P();
    mutex.P();
    w_waiting -= 1;
}
w_active += 1;
mutex.V();
// ... write the file ...
```

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P();
    w_waiting -= 1;
}
w_active += 1;
mutex.V();
// ... write the file ...
```

Implementing “Pass the Baton” (continued)

- 3) When unblocking a process, do not leave the mutual exclusion

```
mutex.P();  
r_active -= 1;  
if( r_active == 0 && w_waiting > 0 )  
    writer.V();  
mutex.V();
```

```
mutex.P();  
r_active -= 1;  
if( r_active == 0 && w_waiting > 0 )  
    writer.V();  
else  
    mutex.V();
```

Implementing “Pass the Baton” (continued)

4) You can only unblock one process at a time

```
// ... write the file ...
mutex.P();
w_active -= 1;
if( r_waiting > 0 )
    for( int i = 0; i < r_waiting; i += 1 )
        reader.V();
else if( w_waiting > 0 )
    writer.V();
mutex.V();
```

```
// ... write the file ...
mutex.P();
w_active -= 1;
if( r_waiting > 0 )
    reader.V();
else if( w_waiting > 0 )
    writer.V();
else
    mutex.V();
```

Implementing “Pass the Baton” (continued)

- 5) When the first unblocked process finishes, have it pass the baton to the next one

```
// READERS
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P();
    mutex.P();
    r_waiting -= 1;
}
r_active += 1;
mutex.V();
// ... read the file ...
```

```
// READERS
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P();
    r_waiting -= 1;
}
r_active += 1;
if( r_waiting > 0 )
    reader.V();
else
    mutex.V();
// ... read the file ...
```


The “Pass the Baton” Solution

```
// READERS
mutex.P();
if( w_active > 0 || w_waiting > 0 ){
    r_waiting += 1;
    mutex.V();
    reader.P();
    r_waiting -= 1;
}
r_active += 1;
if( r_waiting > 0 )
    reader.V();
else
    mutex.V();
// ... read the file ...
mutex.P();
r_active -= 1;
if( r_active == 0 && w_waiting > 0 )
    writer.V();
else
    mutex.V();
```

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P();
    w_waiting -= 1;
}
w_active += 1;
mutex.V();
// ... write the file ...
mutex.P();
w_active -= 1;
if( r_waiting > 0 )
    reader.V();
else if( w_waiting > 0 )
    writer.V();
else
    mutex.V();
```

Pass the Baton: Characteristics

The unblocked processes take over ownership of the mutual exclusion from the unblocking processes

- The unblocked process can access the system state
- The unblocking process must no longer access system state
- Can only unblock one process at a time
- Cohesion is better
- Looks like the P's and V's are not balanced, but they are

Pass the Baton: Characteristics (continued)

Completely invulnerable to “cutting in line”

- 1) One process is reading the file, and a writer arrives

r_active: 1 r_waiting: 0 w_active: 0 w_waiting: 0

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P();
    w_waiting -= 1;
}
w_active += 1;
mutex.V();
// ... write the file ...
```

r_active: 1 r_waiting: 0 w_active: 0 w_waiting: 1

Pass the Baton: Characteristics (continued)

- 2) The reader leaves and tries to unblock the writer

r_active: 1 r_waiting: 0 w_active: 0 w_waiting: 1

```
// WRITERS
mutex.P();
if( r_active > 0 || w_active > 0 ){
    w_waiting += 1;
    mutex.V();
    writer.P(); ← process paused here
    w_waiting -= 1;
}
w_active += 1;
mutex.V();
// ... write the file ...
```

```
// ... read the file ...
mutex.P();
r_active -= 1;
if( r_active == 0 && w_waiting > 0 )
    writer.V();
else
    mutex.V();
```


r_active: 0 r_waiting: 0 w_active: 0 w_waiting: 1

- 3) Newly arriving processes cannot cut in line because they cannot enter the mutual exclusion

Experience

I have used these patterns in my Operating Systems class for several years, and began to describe them more formally for almost two years.

- Students seem to be able to use the patterns easily
- No significant change in grades



Availability

`http://www.cs.rit.edu/~kar/papers` has links to:

- Class handouts describing the patterns and showing their use in several classic synchronization problems
- Descriptions of several problems I have assigned in the past
- More

Thank you!